

CS411ML

April 5, 2025

```
[112]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import statsmodels.api as sm
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LassoCV
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error,
                           mean_absolute_percentage_error
```

```
[92]: df = pd.read_csv('Data for ML - Sheet1.csv')
df.head()
```

```
[92]:      station  year  month  day  avg_wind_speed  max_wind_gust  avg_wind_dir \
0  Belleville  2022      1     1        153.4         22.0       29.6
1  Belleville  2022      1     2        301.9         21.1       14.9
2  Belleville  2022      1     3        190.5         15.1       20.2
3  Belleville  2022      1     4        167.3         24.2       31.9
4  Belleville  2022      1     5        276.9         37.1       32.3

      max_air_temp  avg_air_temp  sol_rad  ...  avg_soil_temp_4in_sod  \
0            86.0        0.01    22.4  ...             48.1
1            71.2        0.02     7.7  ...             41.7
2            59.5        0.02      6  ...             38.4
3            75.0        0.03    24.5  ...             37.9
4            65.6        0.02   12.9  ...             38.0

      max_soil_temp_8in_sod  min_soil_temp_8in_sod  avg_soil_temp_8in_sod  \
0                  50.4                 47.1                49.0
1                  47.1                 41.8                44.2
2                  41.8                 39.7                40.4
3                  39.7                 38.8                39.2
4                  39.7                 39.1                39.3
```

```

max_soil_temp_4in_bare min_soil_temp_4in_bare avg_soil_temp_4in_bare \
0                  53.7                 42.3                47.6
1                  42.3                 35.1                37.5
2                  35.1                 33.5                33.9
3                  33.9                 33.3                33.5
4                  34.0                 33.6                33.8

max_soil_temp_2in_bare min_soil_temp_2in_bare avg_soil_temp_2in_bare
0                  55.2                 38.9                46.1
1                  38.9                 32.2                34.4
2                  32.2                 29.8                31.5
3                  32.1                 30.2                31.5
4                  32.5                 32.0                32.2

```

[5 rows x 26 columns]

```
[93]: soil_columns = [col for col in df.columns if 'soil' in col]

# Determine which soil columns to drop (all except 'avg_soil_temp_8in_sod')
cols_to_drop = [col for col in soil_columns if col != 'avg_soil_temp_8in_sod']
cols_to_drop += ['station', 'year', 'month', 'day']

# Create a tailored DataFrame by dropping the undesired soil columns
df_tailored = df.drop(columns=cols_to_drop)
print(df_tailored.head())
```

	avg_wind_speed	max_wind_gust	avg_wind_dir	max_air_temp	avg_air_temp	\
0	153.4	22.0	29.6	86.0	0.01	
1	301.9	21.1	14.9	71.2	0.02	
2	190.5	15.1	20.2	59.5	0.02	
3	167.3	24.2	31.9	75.0	0.03	
4	276.9	37.1	32.3	65.6	0.02	

	sol_rad	precip	pot_evapot	max_rel_hum	min_rel_hum	avg_soil_temp_8in_sod
0	22.4	33.0	32.9	35.4	35.7	49.0
1	7.7	33.0	33.0	34.8	35.2	44.2
2	6	32.9	32.8	34.2	34.6	40.4
3	24.5	32.8	32.8	34.0	34.1	39.2
4	12.9	32.8	32.7	33.7	33.8	39.3

```
[94]: numeric_cols = df_tailored.select_dtypes(include=['float64', 'int64'])

# Compute the correlation matrix
corr_matrix = numeric_cols.corr()

plt.figure(figsize=(10, 8))
ax = sns.heatmap(corr_matrix,
                  cmap='coolwarm',
```


```

        fmt=".2f",
        linewidths=0.5,
        square=True,
        cbar=True)

# Loop over the data and create text annotations.
for i in range(corr_matrix.shape[0]):
    for j in range(corr_matrix.shape[1]):
        text = f'{corr_matrix.iloc[i, j]:.2f}'
        ax.text(j + 0.5, i + 0.5, text,
                ha="center", va="center", color="black", fontsize=10)

plt.xticks(rotation=45)
plt.yticks(rotation=45)
plt.title("Correlation Matrix Heatmap")
plt.tight_layout()
plt.show()

```



we choose avg_wind_dir, precip, pot_evapot, min_rel_hum as our predictor variables.

```
[95]: if df.index.name == 'date':
    df = df.reset_index()

    # Create a datetime column from 'year', 'month', and 'day'
    df['date'] = pd.to_datetime(df[['year', 'month', 'day']])

    # Sort the DataFrame by date and set 'date' as the index
    df = df.sort_values('date')
    df.set_index('date', inplace=True)


    # Convert to daily frequency; if some days are missing, they become NaN
    df = df.asfreq('D')

    # Remove rows with missing values in 'avg_soil_temp_8in_sod'
    df = df.dropna(subset=['avg_soil_temp_8in_sod'])

    # Remove any duplicate dates (if they exist)
    df = df[~df.index.duplicated(keep='first')]

    # Plot the time series for avg_soil_temp_8in_sod
    plt.figure(figsize=(12, 6))
    sns.lineplot(x=df.index, y=df['avg_soil_temp_8in_sod'], color='tab:red', alpha=0.7)
    plt.xlabel("Date")
    plt.ylabel("Soil Temperature (8in sod)")
    plt.title("Daily Soil Temperature Time Series Plot")
    plt.xticks(rotation=45)
    plt.grid(True)
    plt.tight_layout()
    plt.show()
```



```
C:\Users\54206\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\54206\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```



```
[96]: y = df['avg_soil_temp_8in_sod']

# 5. Plot ACF
plt.figure(figsize=(12, 4))
sm.graphics.tsa.plot_acf(y, lags=30, ax=plt.gca())
plt.title("ACF of avg_soil_temp_8in_sod")
plt.xlabel("Lags")
plt.ylabel("Autocorrelation")
plt.tight_layout()
plt.show()

# 6. Plot PACF
plt.figure(figsize=(12, 4))
sm.graphics.tsa.plot_pacf(y, lags=30, ax=plt.gca(), method='ywm')
plt.title("PACF of avg_soil_temp_8in_sod")
plt.xlabel("Lags")
plt.ylabel("Partial Autocorrelation")
plt.tight_layout()
plt.show()
```



The best lag at 2

```
[97]: best_lag = 2
for lag in range(1, best_lag + 1):
    df[f'Tlag_{lag}'] = y.shift(lag)
df.dropna(inplace=True)
```

```
[102]: X = df[['avg_wind_dir',
             'precip',
             'pot_evapot',
             'min_rel_hum',
             'Tlag_1',
             'Tlag_2']]
Y = df[['avg_soil_temp_8in_sod']]
X = sm.add_constant(X)
y = y.loc[X.index]
```

```
[107]: # Split the data into training (80%) and testing (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)
```

```

# Generate polynomial features (degree 2 includes quadratic and interaction terms)
poly = PolynomialFeatures(degree=2, include_bias=False)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

# Fit a linear regression model on the polynomial features
model = LinearRegression()
model.fit(X_train_poly, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test_poly)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)

# Print the evaluation metrics with high precision
print("Evaluation Metrics:")
print(f"MSE: {mse:.3f}")
print(f"RMSE: {rmse:.3f}")
print(f"MAE: {mae:.3f}")
print(f"MAPE: {mape:.3f}")

```

Evaluation Metrics:

MSE: 1.191
 RMSE: 1.091
 MAE: 0.710
 MAPE: 0.015

```

[111]: lasso_cv = LassoCV(cv=5, random_state=42)
lasso_cv.fit(X_train_poly, y_train)

print("Best alpha selected:", lasso_cv.alpha_)
y_pred = lasso_cv.predict(X_test_poly)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)

# Print the evaluation metrics with high precision
print("Evaluation Metrics:")
print(f"MSE: {mse:.3f}")

```

```

print(f"RMSE: {rmse:.3f}")
print(f"MAE: {mae:.3f}")
print(f"MAPE: {mape:.3f}")

```

Best alpha selected: 43.55349976779514
Evaluation Metrics:
MSE: 3.470
RMSE: 1.863
MAE: 1.566
MAPE: 0.034

```

[114]: gb_model = GradientBoostingRegressor(random_state=42)
gb_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = gb_model.predict(X_test)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)

# Print the evaluation metrics with high precision
print("Evaluation Metrics:")
print(f"MSE: {mse:.3f}")
print(f"RMSE: {rmse:.3f}")
print(f"MAE: {mae:.3f}")
print(f"MAPE: {mape:.3f}")

```

Evaluation Metrics:
MSE: 1.488
RMSE: 1.220
MAE: 0.860
MAPE: 0.017

Model	MSE	RMSE	MAE	MAPE
Polynomial Regression or	1.191	1.091	0.710	0.015
Lasso Polynomial Regression	3.470	1.863	1.566	0.034

```

[115]: # Define predictors and target variable from your DataFrame
X = df[['avg_wind_dir', 'precip', 'pot_evapot', 'min_rel_hum', 'Tlag_1',
        'Tlag_2']]
y = df['avg_soil_temp_8in_sod']

# Generate quadratic (degree 2) polynomial features
poly = PolynomialFeatures(degree=2, include_bias=False)

```

```

X_poly = poly.fit_transform(X)

# Fit the polynomial regression model on the entire dataset
model = LinearRegression()
model.fit(X_poly, y)

# Predict the soil temperature for the entire year
df['predicted_soil_temp'] = model.predict(X_poly)

# Function to map month to season
def month_to_season(month):
    if month in [12, 1, 2]:
        return 'Winter'
    elif month in [3, 4, 5]:
        return 'Spring'
    elif month in [6, 7, 8]:
        return 'Summer'
    elif month in [9, 10, 11]:
        return 'Fall'

# Create a 'season' column by mapping the month from the datetime index
df['season'] = df.index.month.map(month_to_season)

# Group predictions by season and compute the average predicted soil temperature
seasonal_avg = df.groupby('season')['predicted_soil_temp'].mean()

print("Seasonal Average Predicted Soil Temperature:")
print(seasonal_avg)

```

Seasonal Average Predicted Soil Temperature:

Season	Predicted Soil Temperature
Fall	60.757891
Spring	54.842481
Summer	76.129921
Winter	38.160734

Name: predicted_soil_temp, dtype: float64